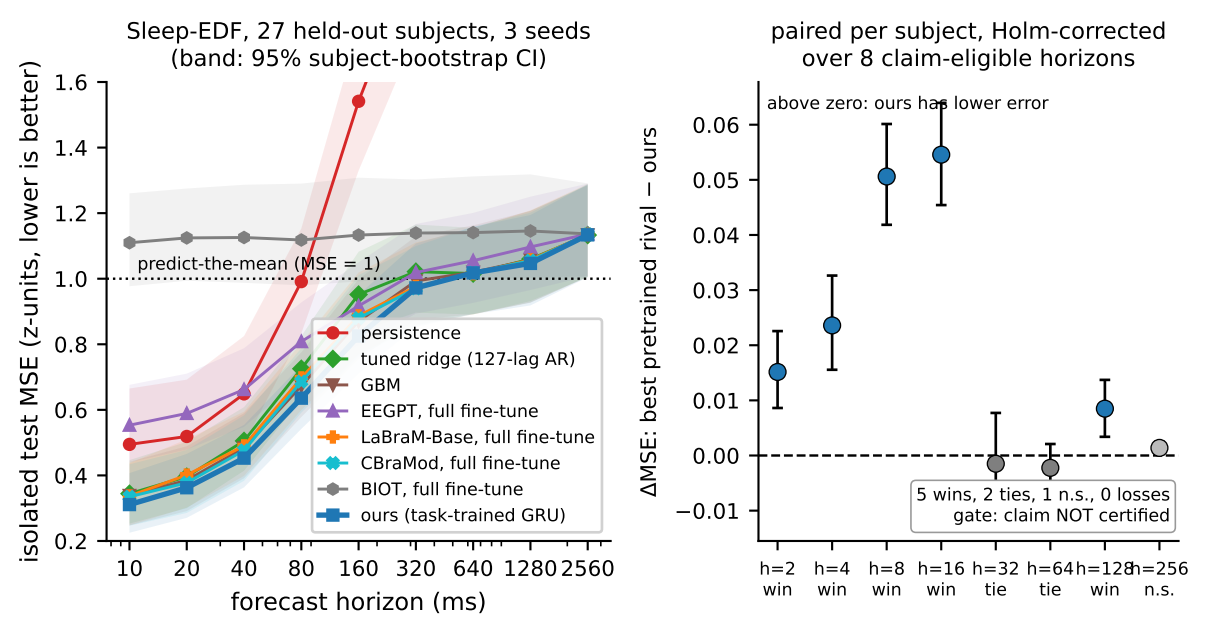
Sleep-EDF, 27 held-out subjects, 3 seeds
(band: 95% subject-bootstrap CI)
isolated test MSE (z-units, lower is better)
forecast horizon (ms)
predict-the-mean (MSE = 1)
persistence
tuned ridge (127-lag AR)
GBM
EEGPT, full fine-tune
LaBraM-Base, full fine-tune
CBraMod, full fine-tune
BIOT, full fine-tune
ours (task-trained GRU)
paired per subject, Holm-corrected
over 8 claim-eligible horizons
above zero: ours has lower error
ΔMSE: best pretrained rival − ours
5 wins, 2 ties, 1 n.s., 0 losses
gate: claim NOT certified
h=2 win
h=4 win
h=8 win
h=16 win
h=32 tie
h=64 tie
h=128 win
h=256 n.s.